

Lecture 22: B Trees I

2023/10/12


詹博华 (中国科学院软件研究所)

B Trees

- Another (balanced) search tree data structure. Aim to minimize disk I/O operations for trees stored on a disk.
- Applications in database systems.
- Large branching factor, compared to 2 for binary search trees.

Properties of disk drive

- Read/write time significantly longer than main memory.
- Disk access not just one item but pages at a time, with each page 2^{11} to 2^{14} bytes (2kb – 16kb) in length.
- Design algorithms to minimize number of disk I/O (rather than total number of computation steps).



Aside: big-O notation: What is an operation?

- We often say an algorithm is $O(f(n))$, meaning the number of operations (or computation steps) grows as at most $c \cdot f(n)$. This depends on what counts as an operation.
- For example, if an algorithm involve multiplication of integers. We may choose to consider one multiplication $x \cdot y$ as one step, or as several steps, the number of which depends on the size of x and y (e.g. $O(\log(x) \cdot \log(y))$ steps).
- In our case, we only consider read/write from disk as operations, and ignore the other computations (this can be added back in a more refined analysis).

Properties of B Trees

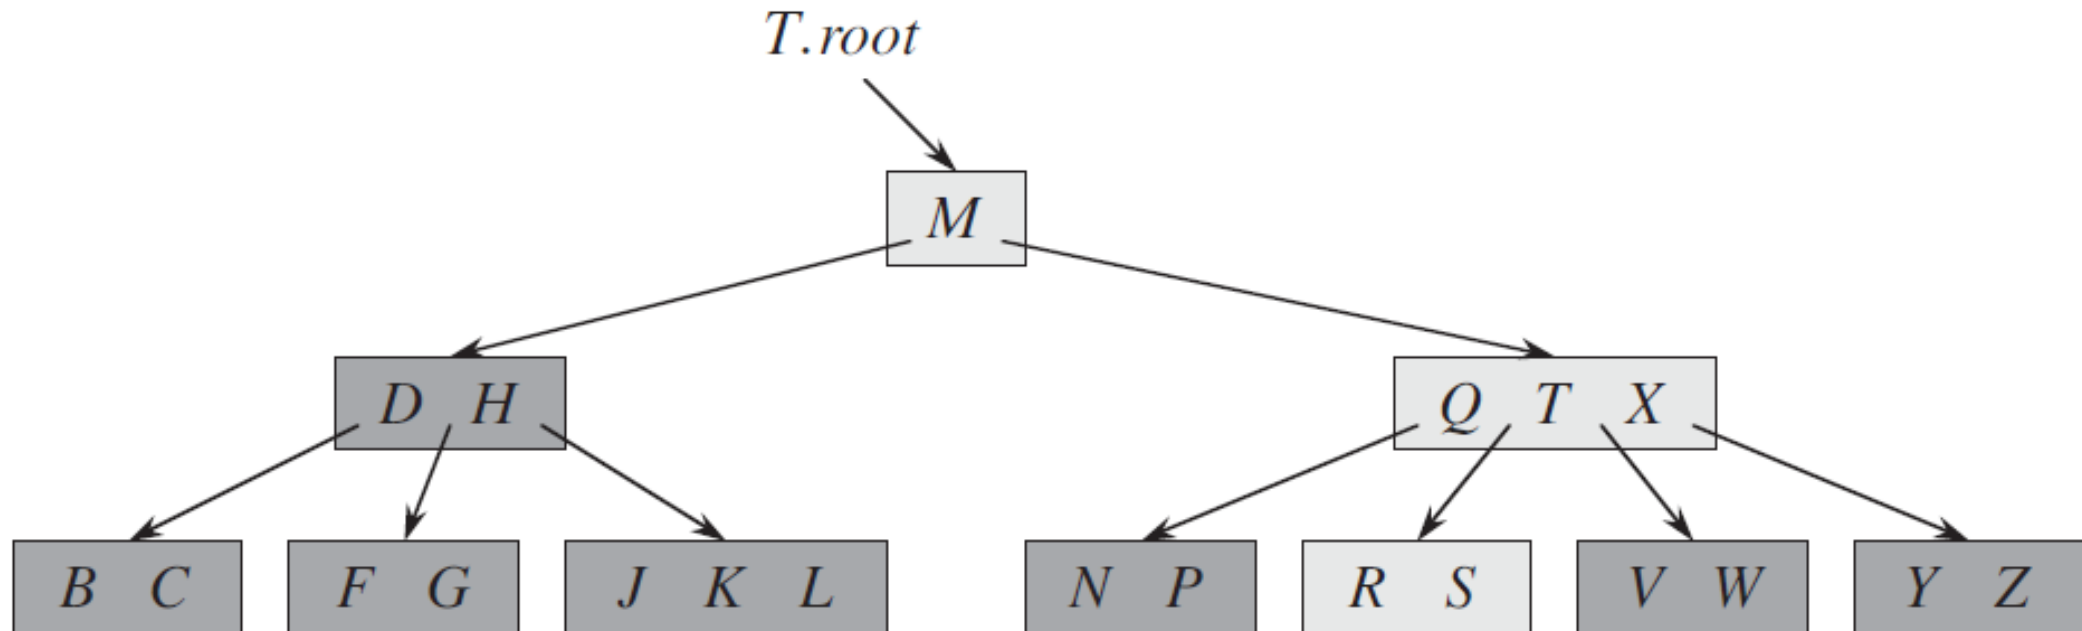
- Each node x stores the number of keys $x.n$, and a list of keys

$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_n.$$


- Each node is either internal node or a leaf. Each internal node contains $x.n + 1$ pointers to its children.
- The keys separate the range stored in each subtree.
- All leaves have the same depth h .
- Given a minimum degree t , each node other than the root must have at least $t - 1$ keys. Each node contains at most $2t - 1$ keys.

Example of B Trees

- B Tree with $h = 3$ and $t = 2$ (minimum 1, maximum 3 entries per node).



Height vs. size of B Trees



depth	number of nodes
0	1
1	2
2	$2t$
3	$2t^2$

Height vs. size of B Trees: Computation

- Lower bound on size in term of height:

$$n \geq 1 + (t - 1) \sum_{i=1}^h 2t^{i-1} = 1 + 2(t - 1) \frac{t^h - 1}{t - 1} = 2t^h - 1$$

- This implies:

$$h \leq \log_t \frac{n + 1}{2}$$

Search in B-Trees

- Standard procedure yields $O(h) = O(\log_t n)$ algorithm (in the number of read/write of blocks).
- Note however the while loop on line 2-3, giving an extra factor of t if we also consider these computations.

B-TREE-SEARCH(x, k)


```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

Insertion in B-Trees: Naïve solution

- Insert as in ordinary search tree.
- Split nodes that become too large in the process.
- Require backing-up the tree – extra constant factor in number of disk I/O's.

Splitting of nodes

- Starting from a node that is full ($2t - 1$ nodes).
- End with two nodes that have minimum size ($t - 1$ nodes).
- Example for $t = 4$:





Insertion in a single pass

- During the search for location of insertion, split any full node that is encountered.
- This ensures that insertion is always made to non-full nodes.



Insertion: example ($t = 3$)

- Insert B, resulting in a full leaf.



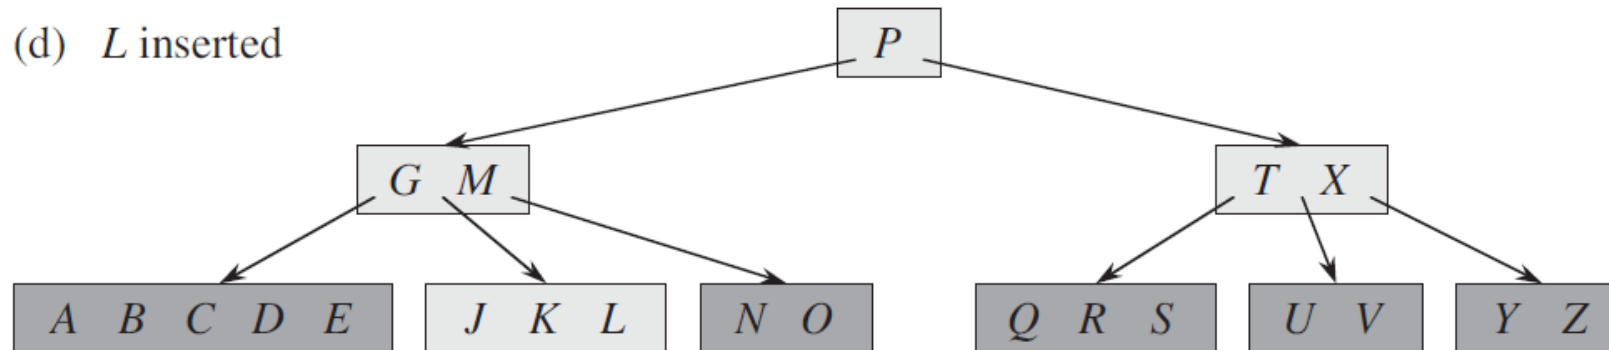

Insertion: example

- To insert Q , first split the node $\langle R S T U V \rangle$, moving T to the root.




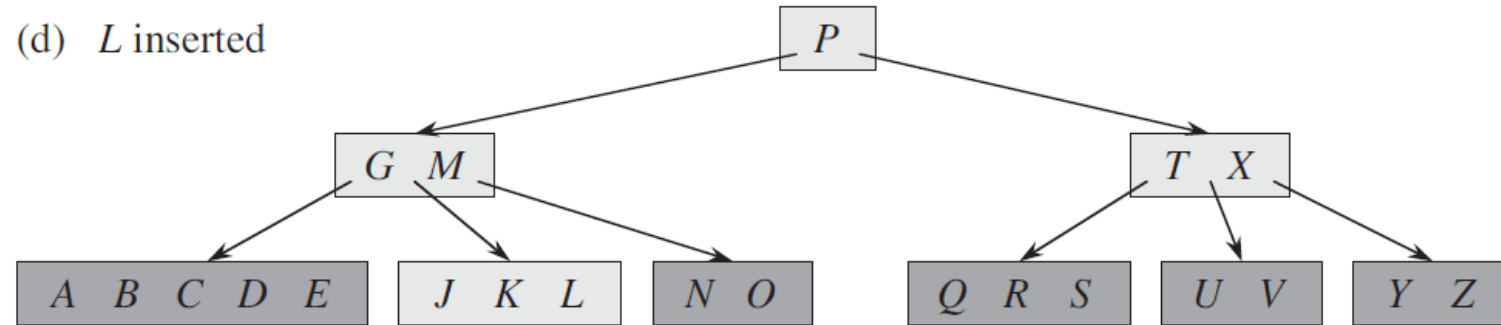
Insertion: example

- Before insertion L , first split root node $\langle G \ M \ P \ T \ X \rangle$.



Insertion: example

- Before inserting F , first split node $\langle A\ B\ C\ D\ E\rangle$.



Insertion: implementation

- Line 2-8: when a full node is encountered, split node.
- Line 9, 10: insert into non-full nodes.

B-TREE-INSERT(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

Insertion: non-full case

- Line 2-8: leaf case.
- Line 9-17: non-leaf case.

B-TREE-INSERT-NONFULL(x, k)

```
1   $i = x.n$ 
2  if  $x.leaf$ 
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$ 
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$ 
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```